

On integrating Software-Defined Networking within existing routing systems



Laurent Vanbever

Princeton University

Stanford University

November, 13 2013

On integrating Software-Defined Networking within existing routing systems



- 1 **SDN-controlled routers**
don't trash, recycle
- 2 **SDN-controlled IGP**
fine-grained traffic-engineering
- 3 **SDN-controlled BGP**
inter domain bonanza

On integrating Software-Defined Networking within existing routing systems



1 **SDN-controlled routers**
don't trash, recycle

SDN-controlled IGP
fine-grained traffic-engineering

SDN-controlled BGP
inter domain bonanza

Today's networks are managed indirectly

Given network-wide forwarding requirements

Traffic from i to j should flow along path $P1$

Traffic from k to l should flow along path $P2$

...

operators' job

Configure each equipment such that
they compute (locally) compatible forwarding entries

Today's networks are managed indirectly, device-by-device

Given network-wide forwarding requirements

Traffic from i to j should flow along path $P1$

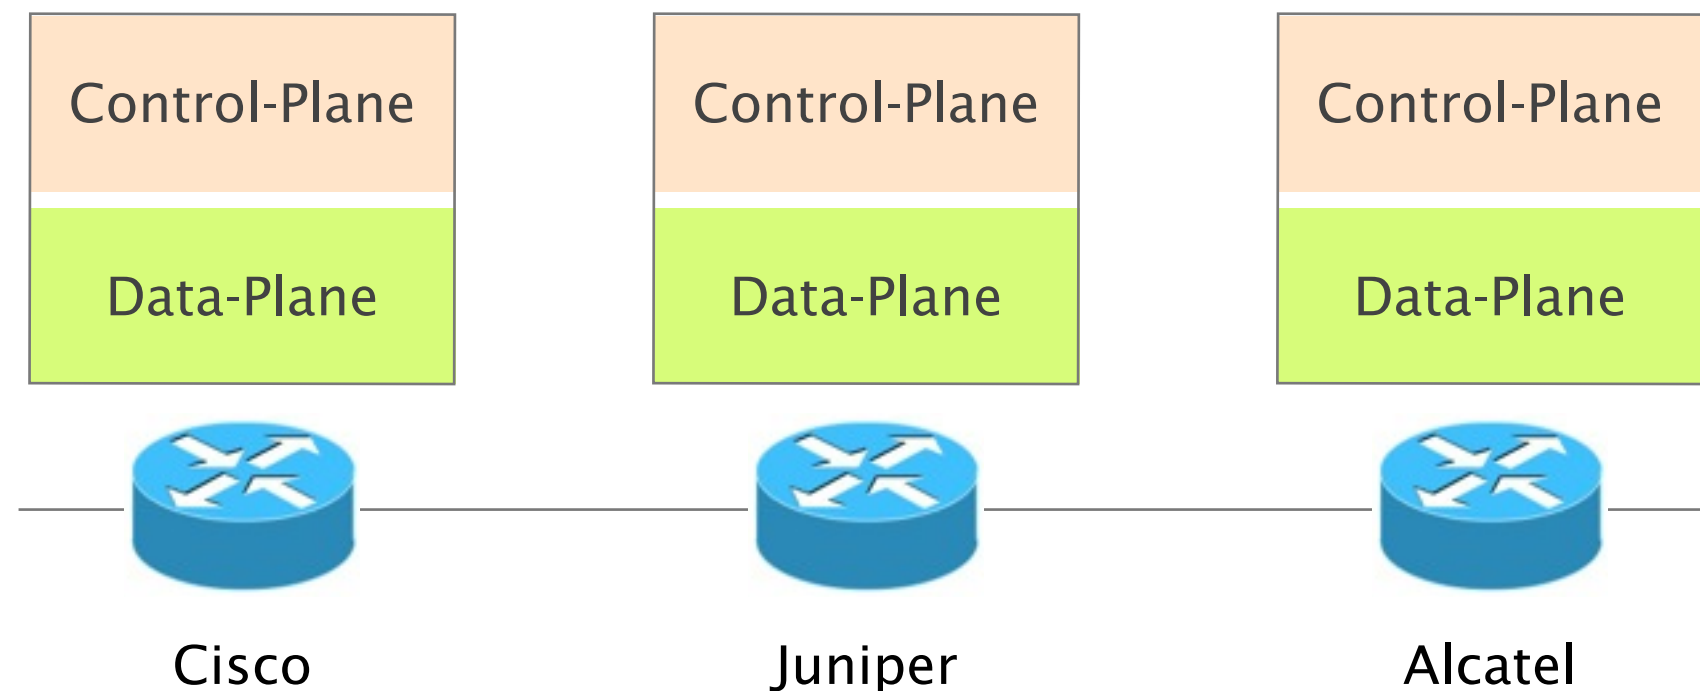
Traffic from k to l should flow along path $P2$

...

operators' job

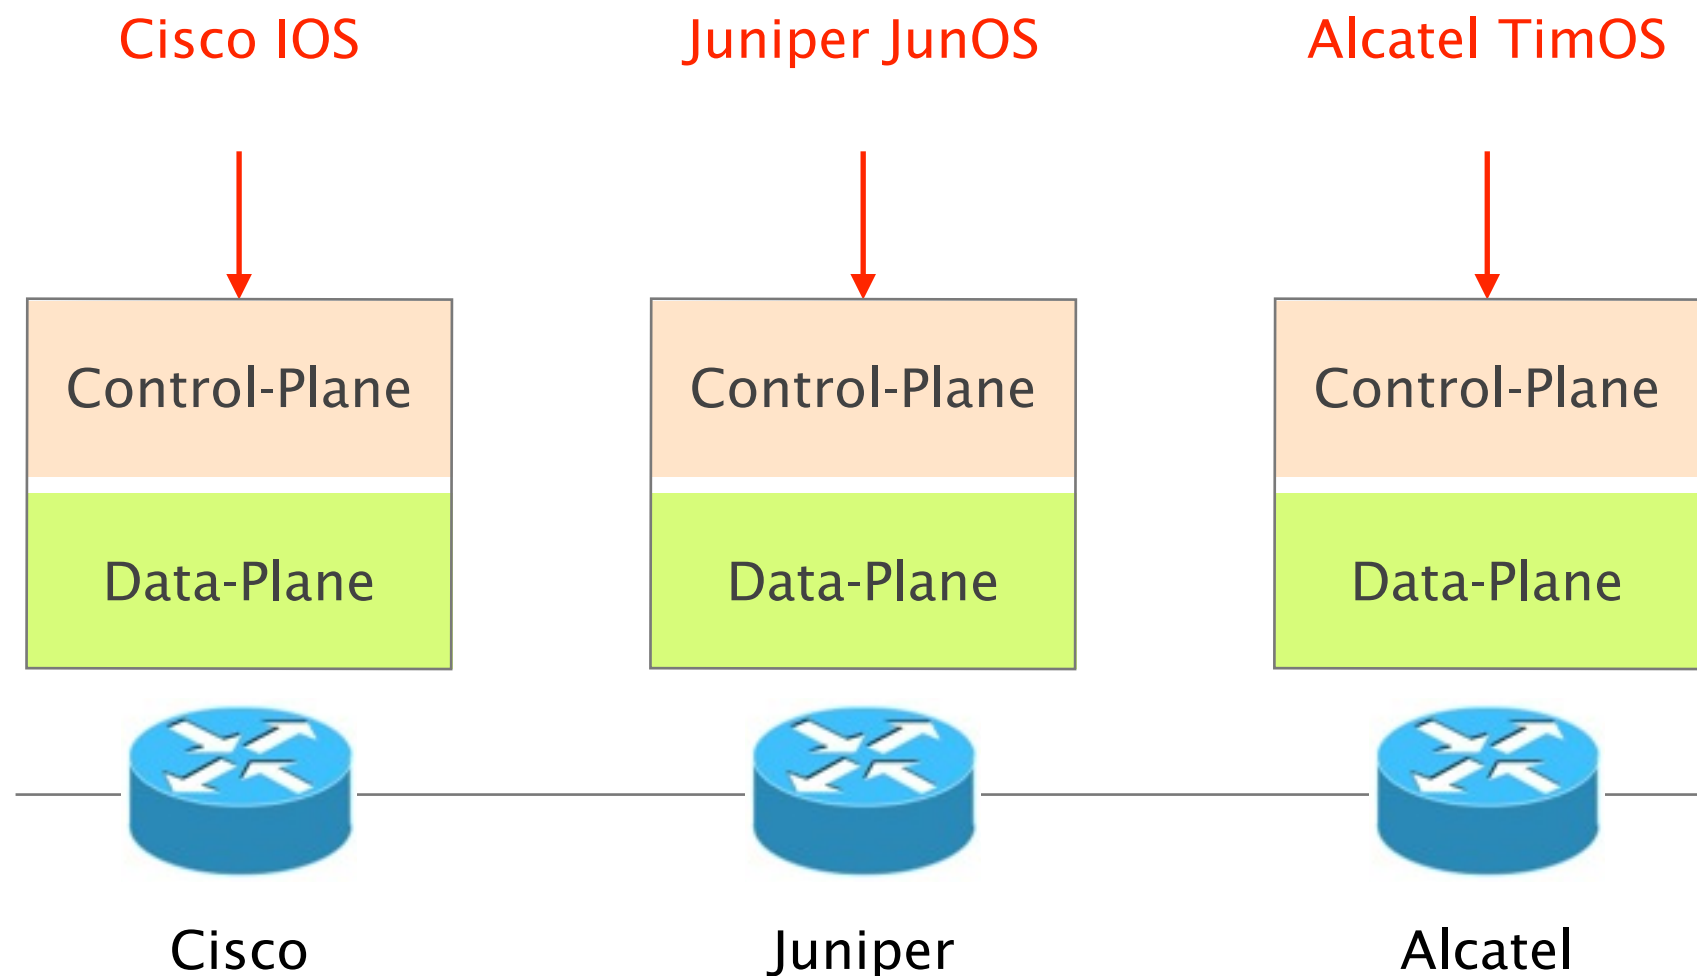
Configure each equipment such that
they compute (locally) compatible forwarding entries

Today's networks are managed indirectly,
device-by-device, using arcane configuration languages



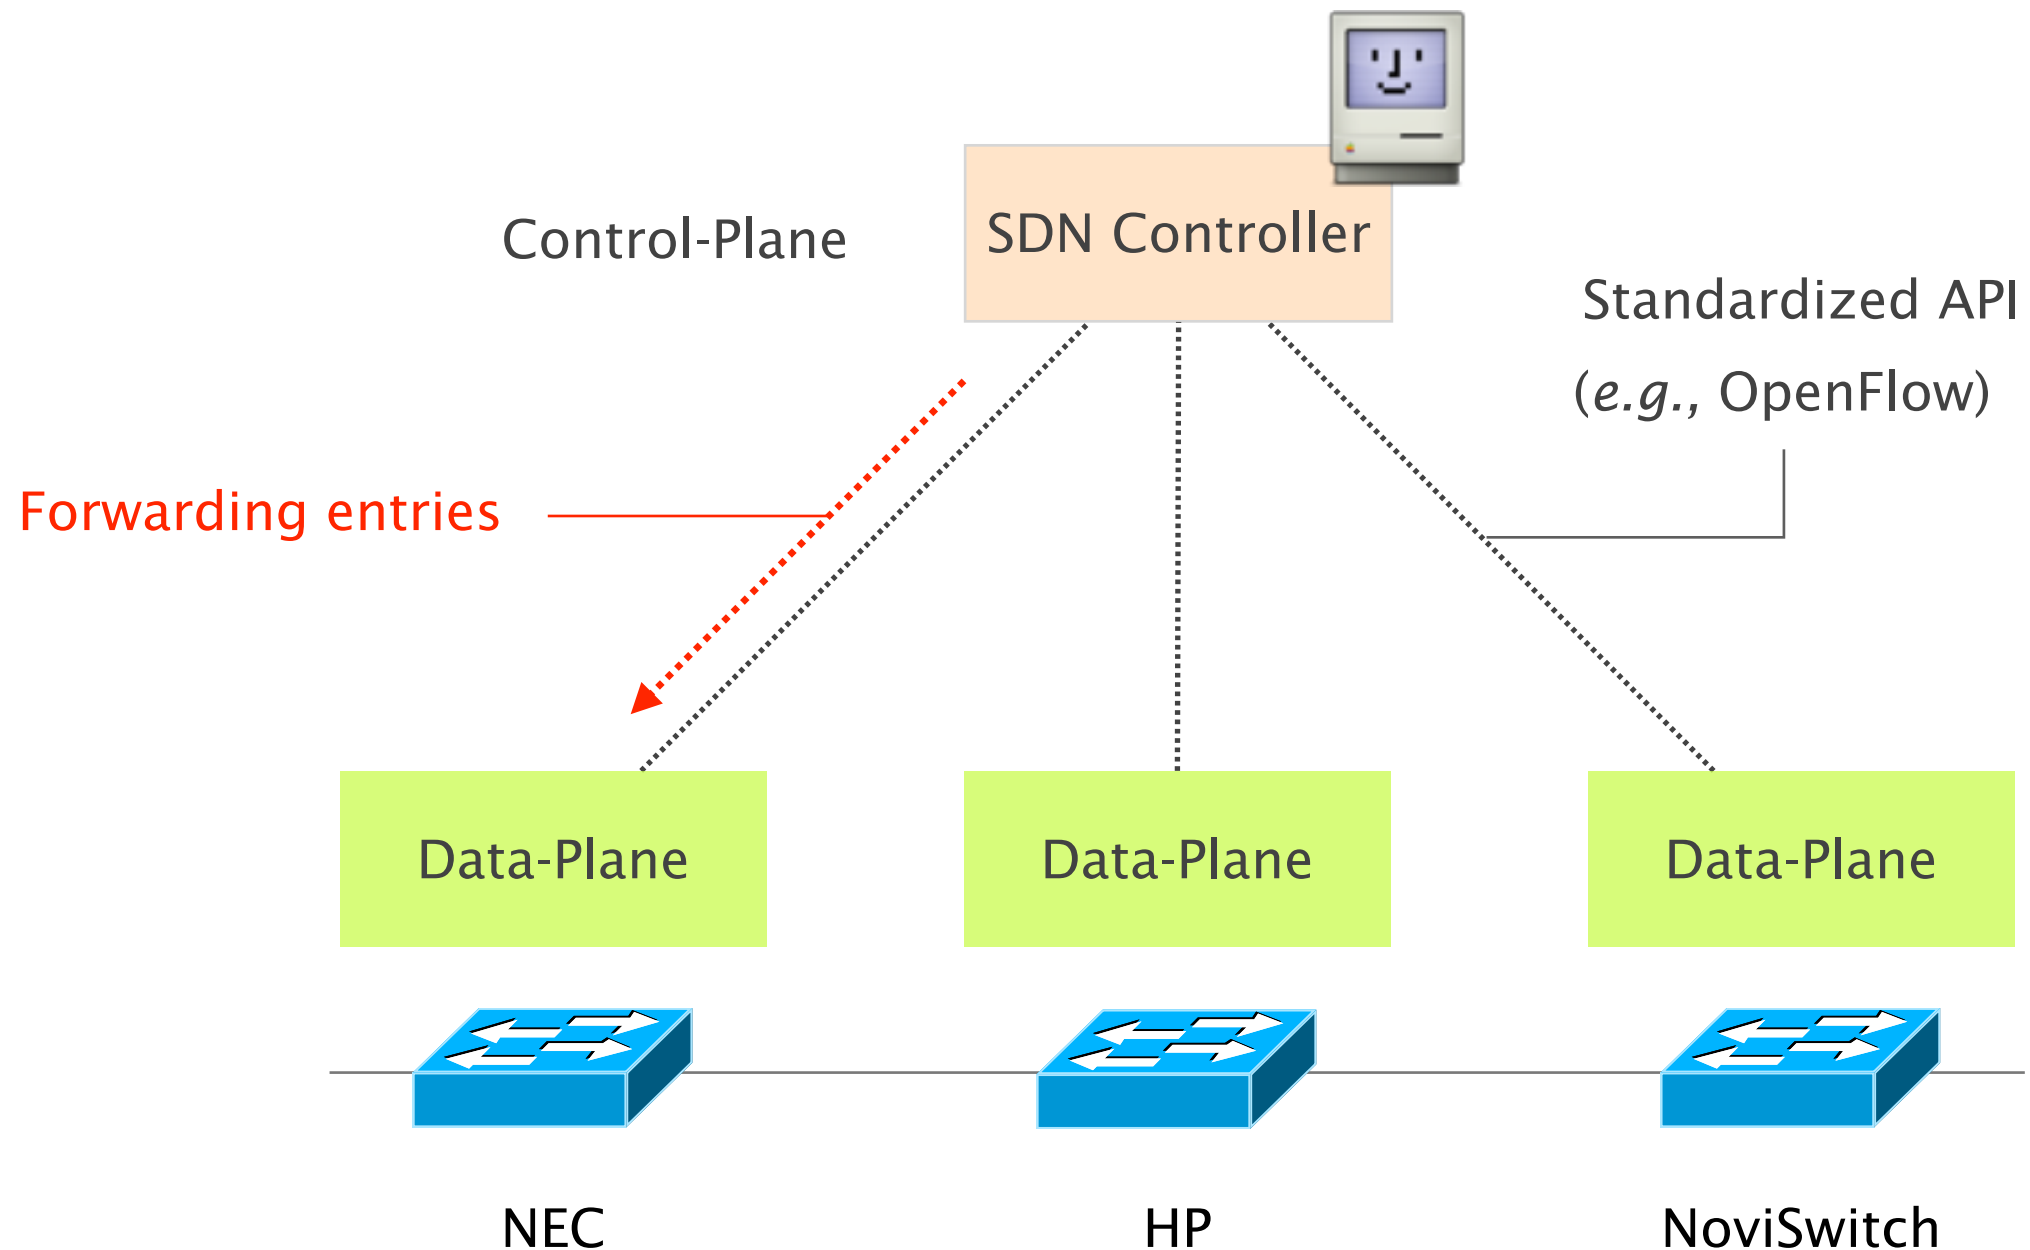
Today's networks are managed indirectly,
device-by-device, using arcane configuration languages

configuration
"languages"



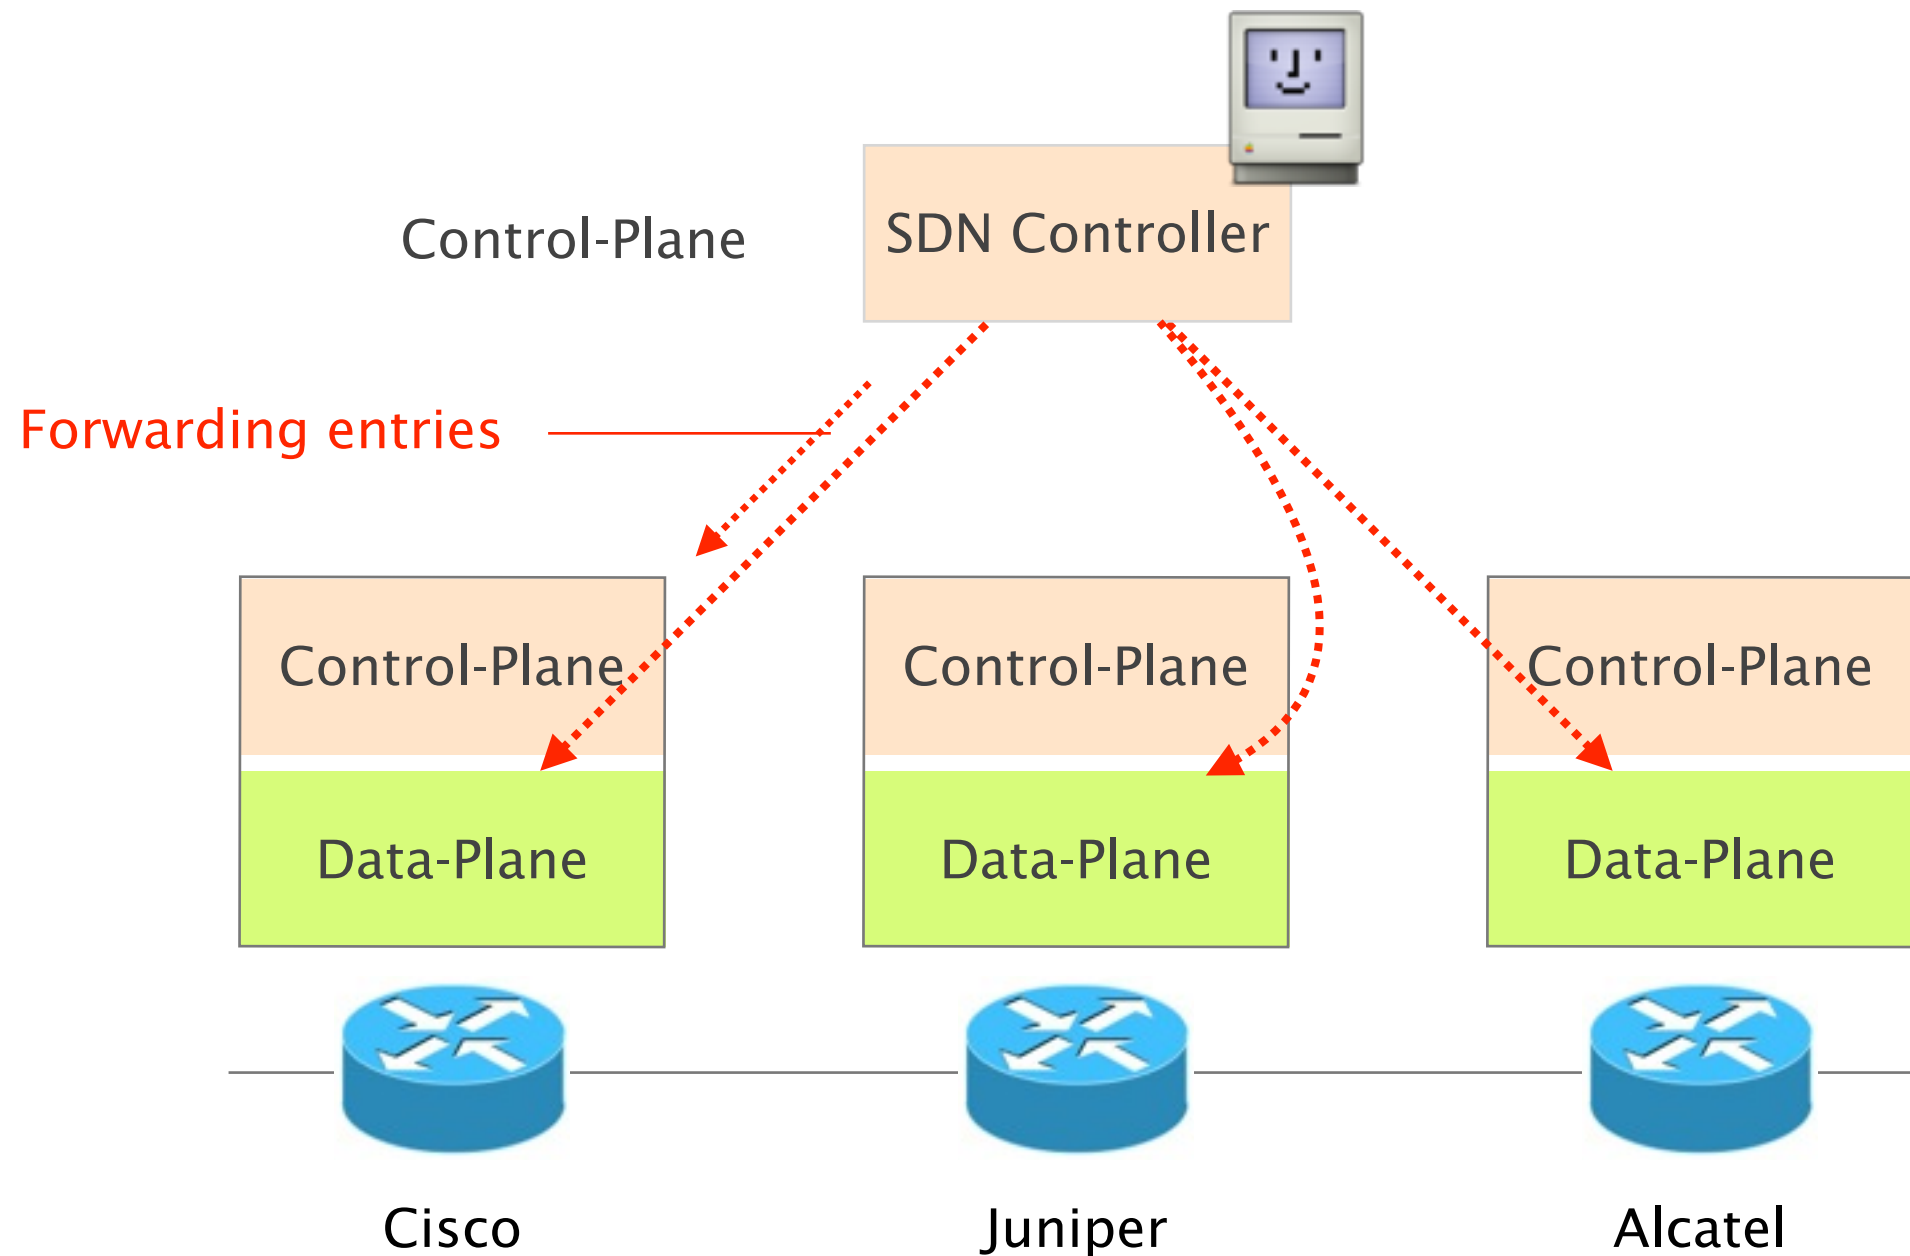
In contrast, SDN simplifies network management...

...by directly programming forwarding entries,
using a logically-centralized controller and an open API

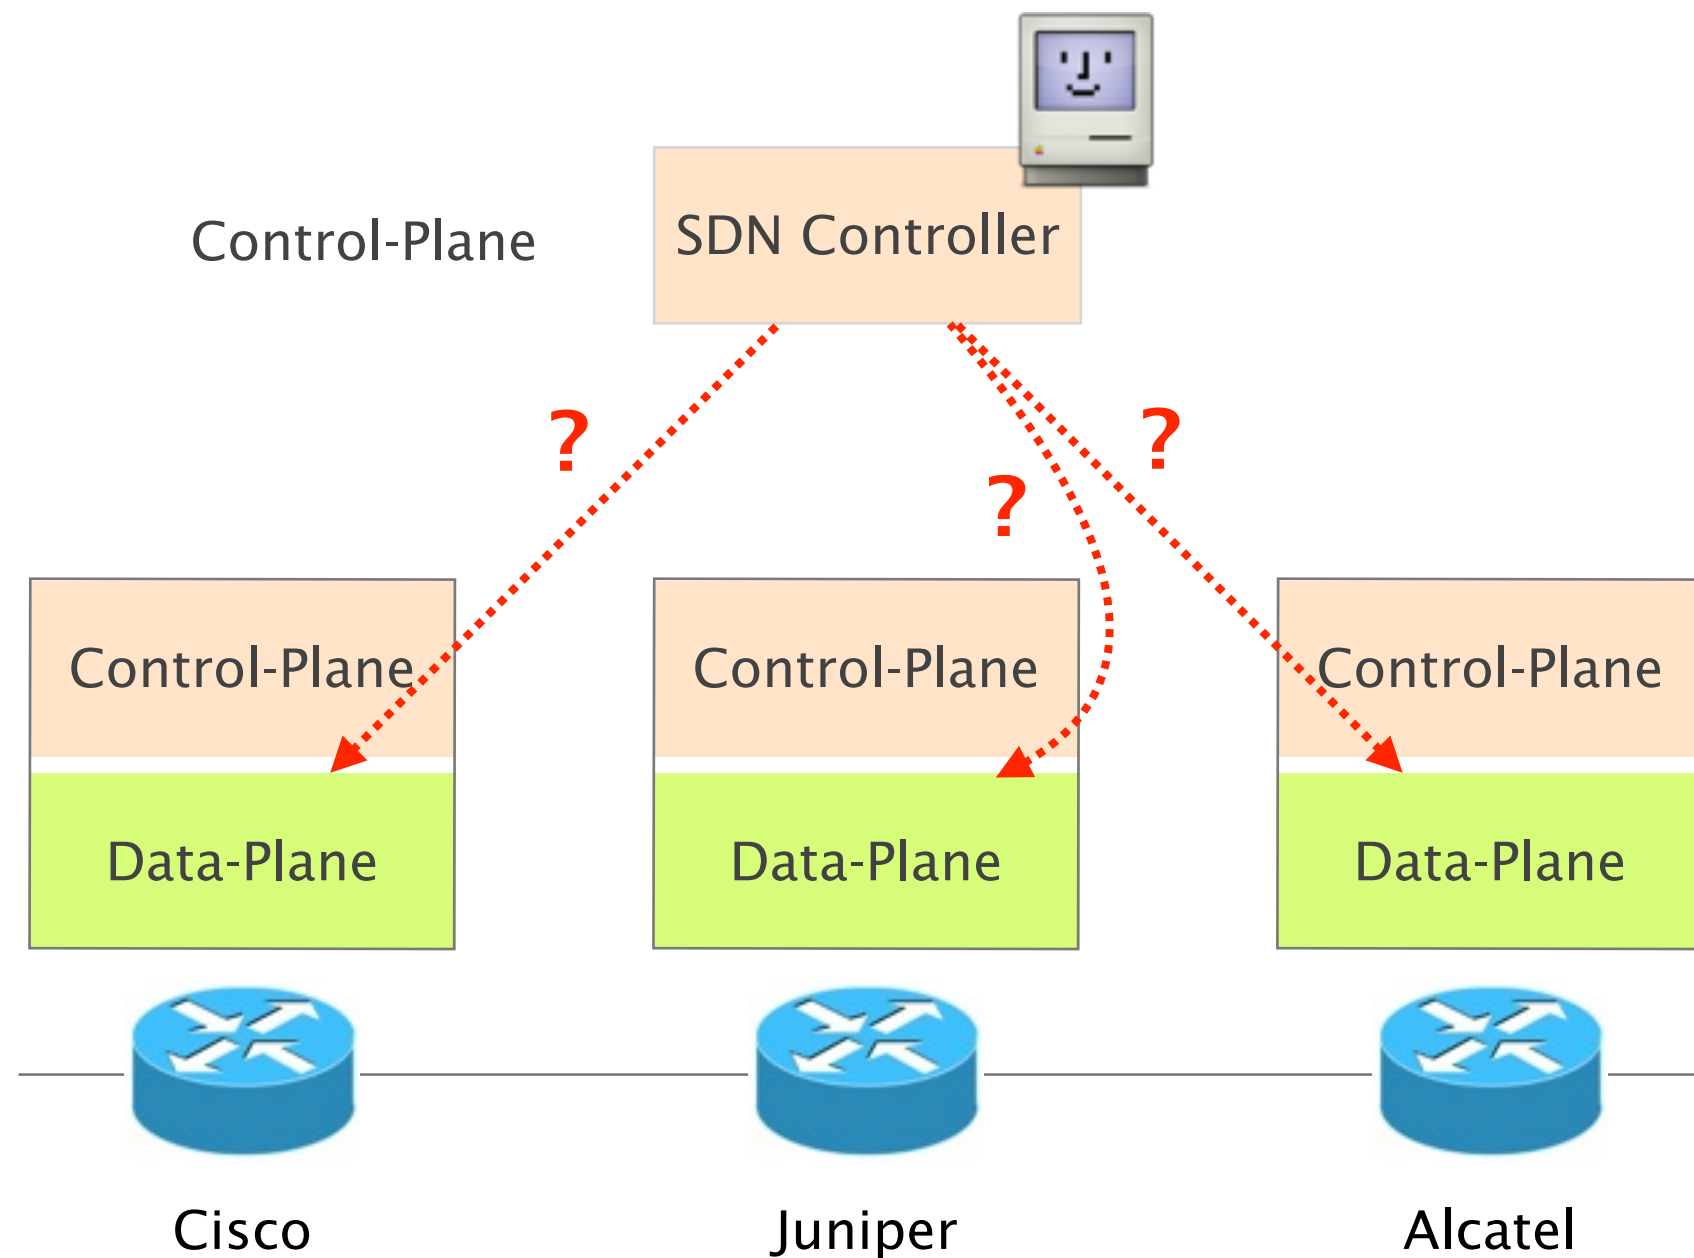


The bad news is that SDN requires compatible devices...

Wouldn't it be great to manage
an *existing* network “à la SDN”?



To do that, we need an open API to
program forwarding entries in a router

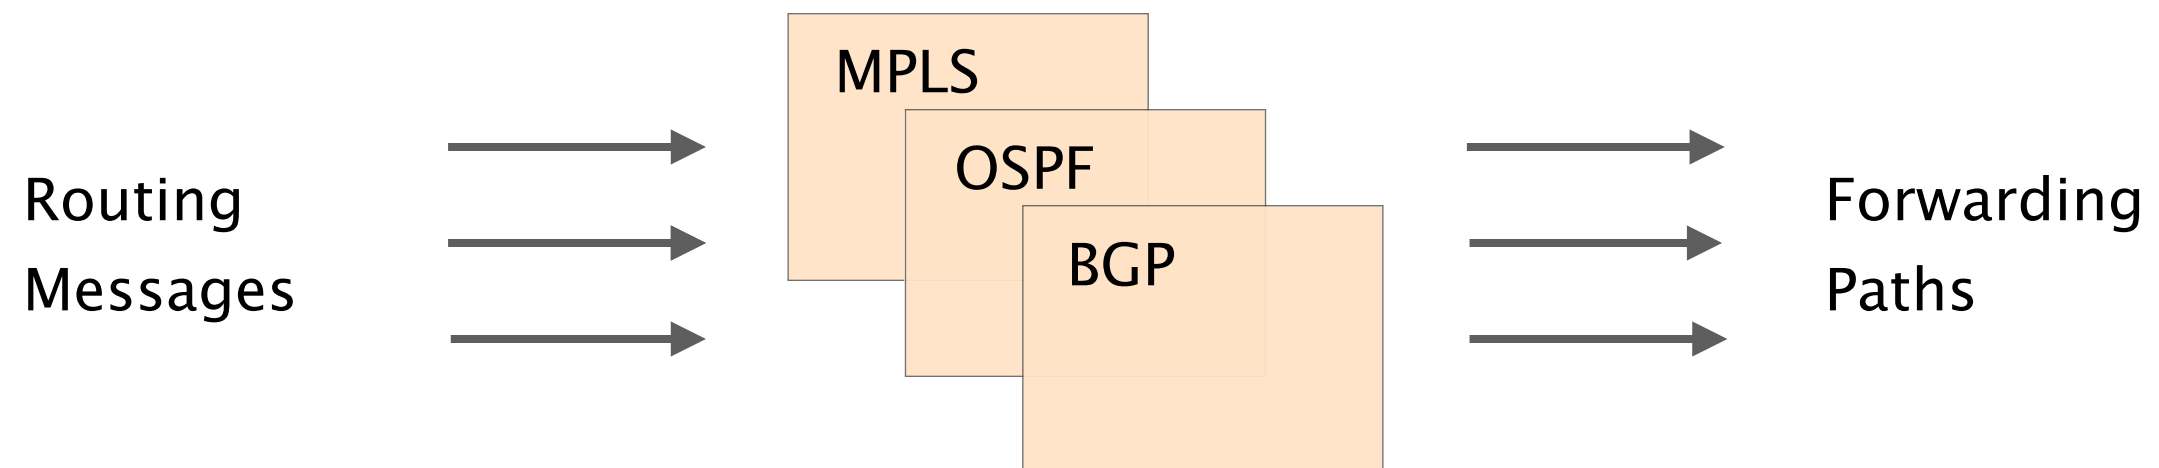


Routing protocols are good candidates to act as API


Routing protocols

- messages are standardized
all boxes must speak the same language
- behaviors are well-defined and understood
e.g., shortest-path routing
- implementations are widely available
a vast majority (if not all) Cisco devices supports OSPF

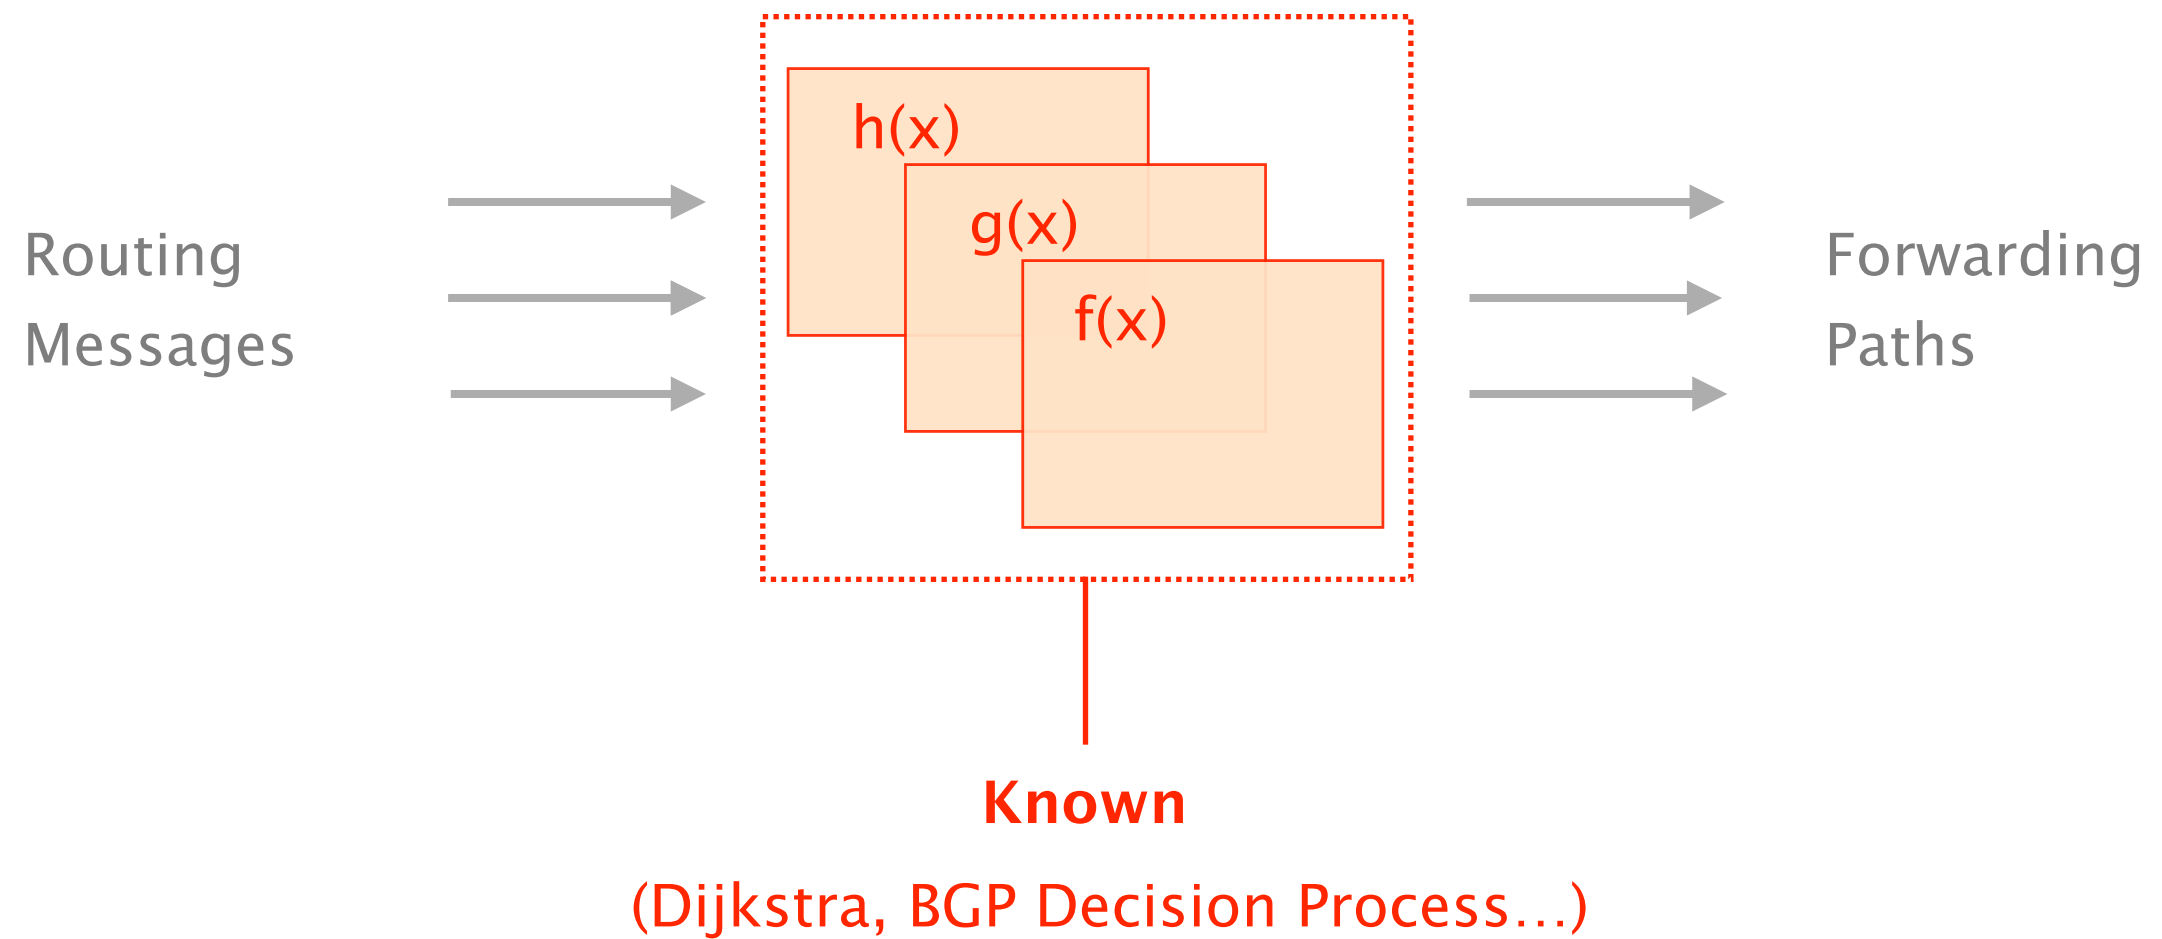
A routing protocol takes routing messages as input and computes forwarding paths as output



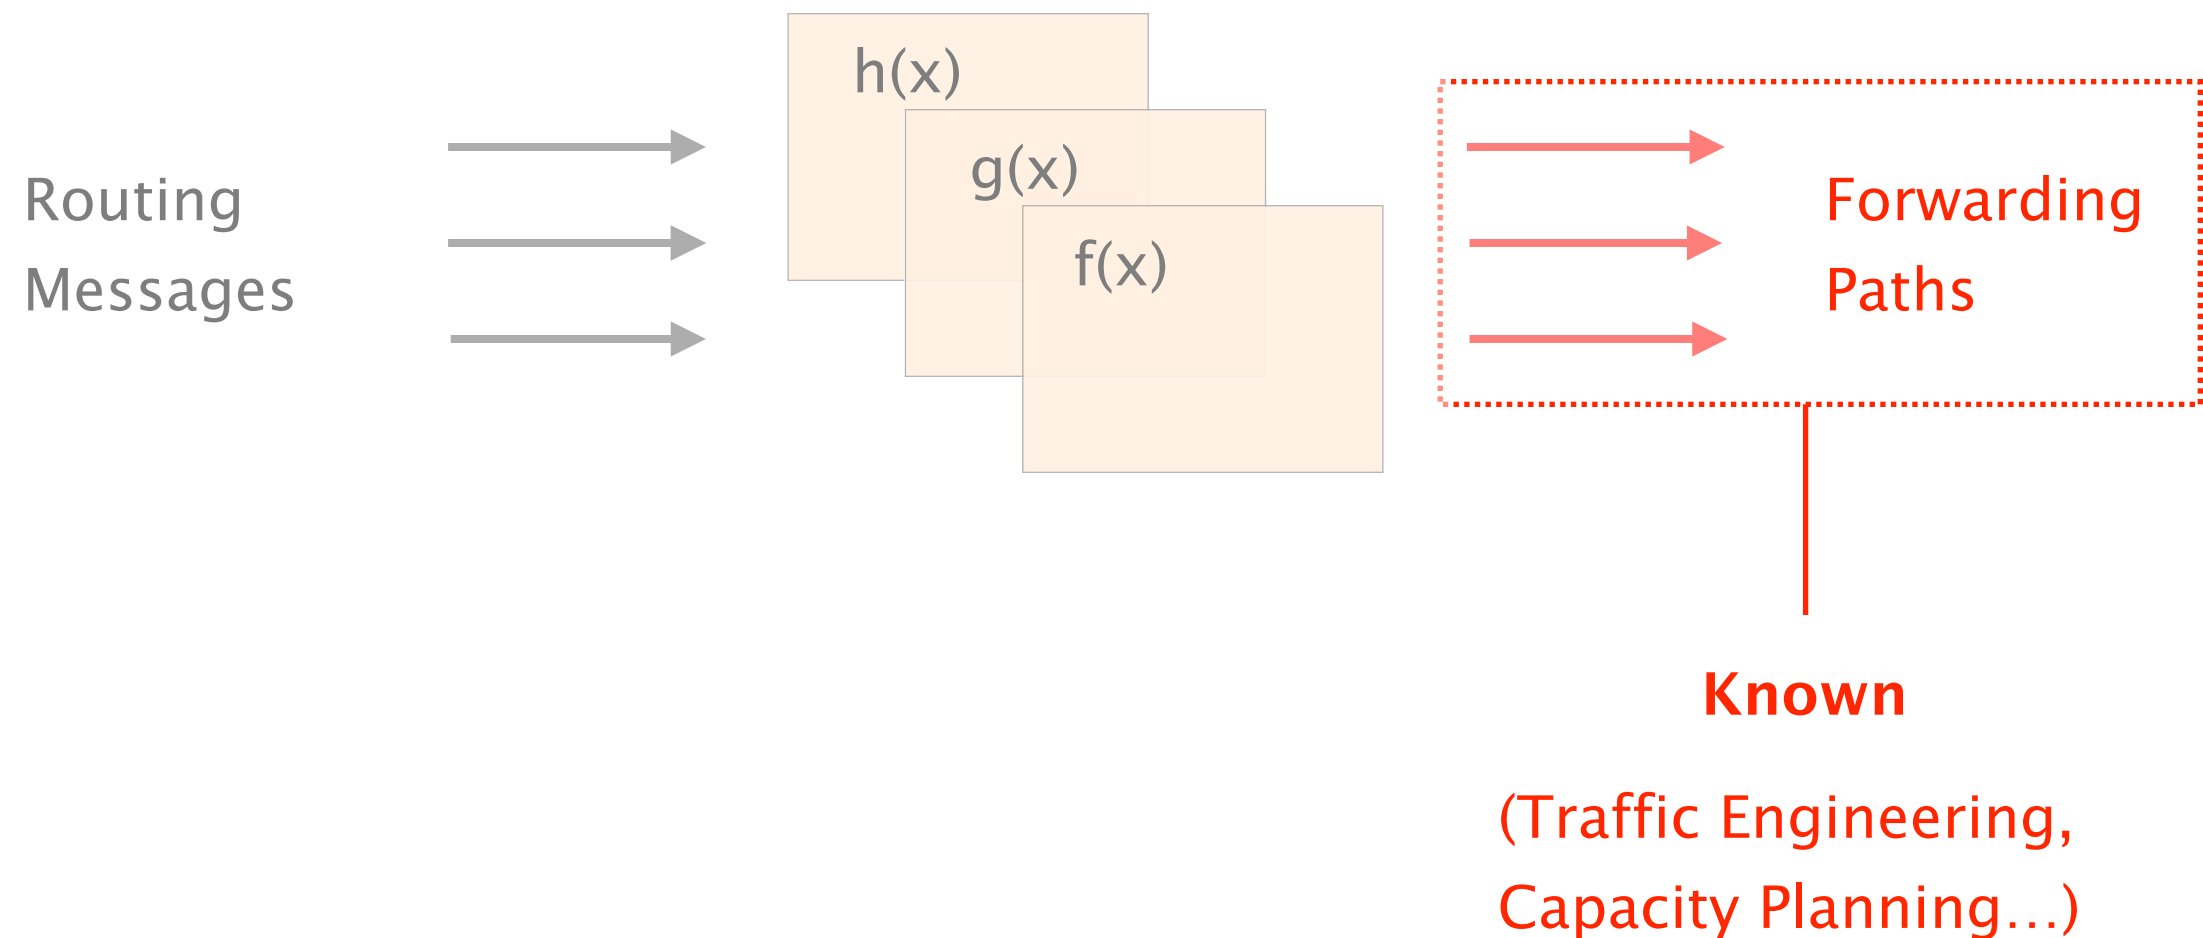
A routing protocol is thus a function
from input messages to forwarding paths



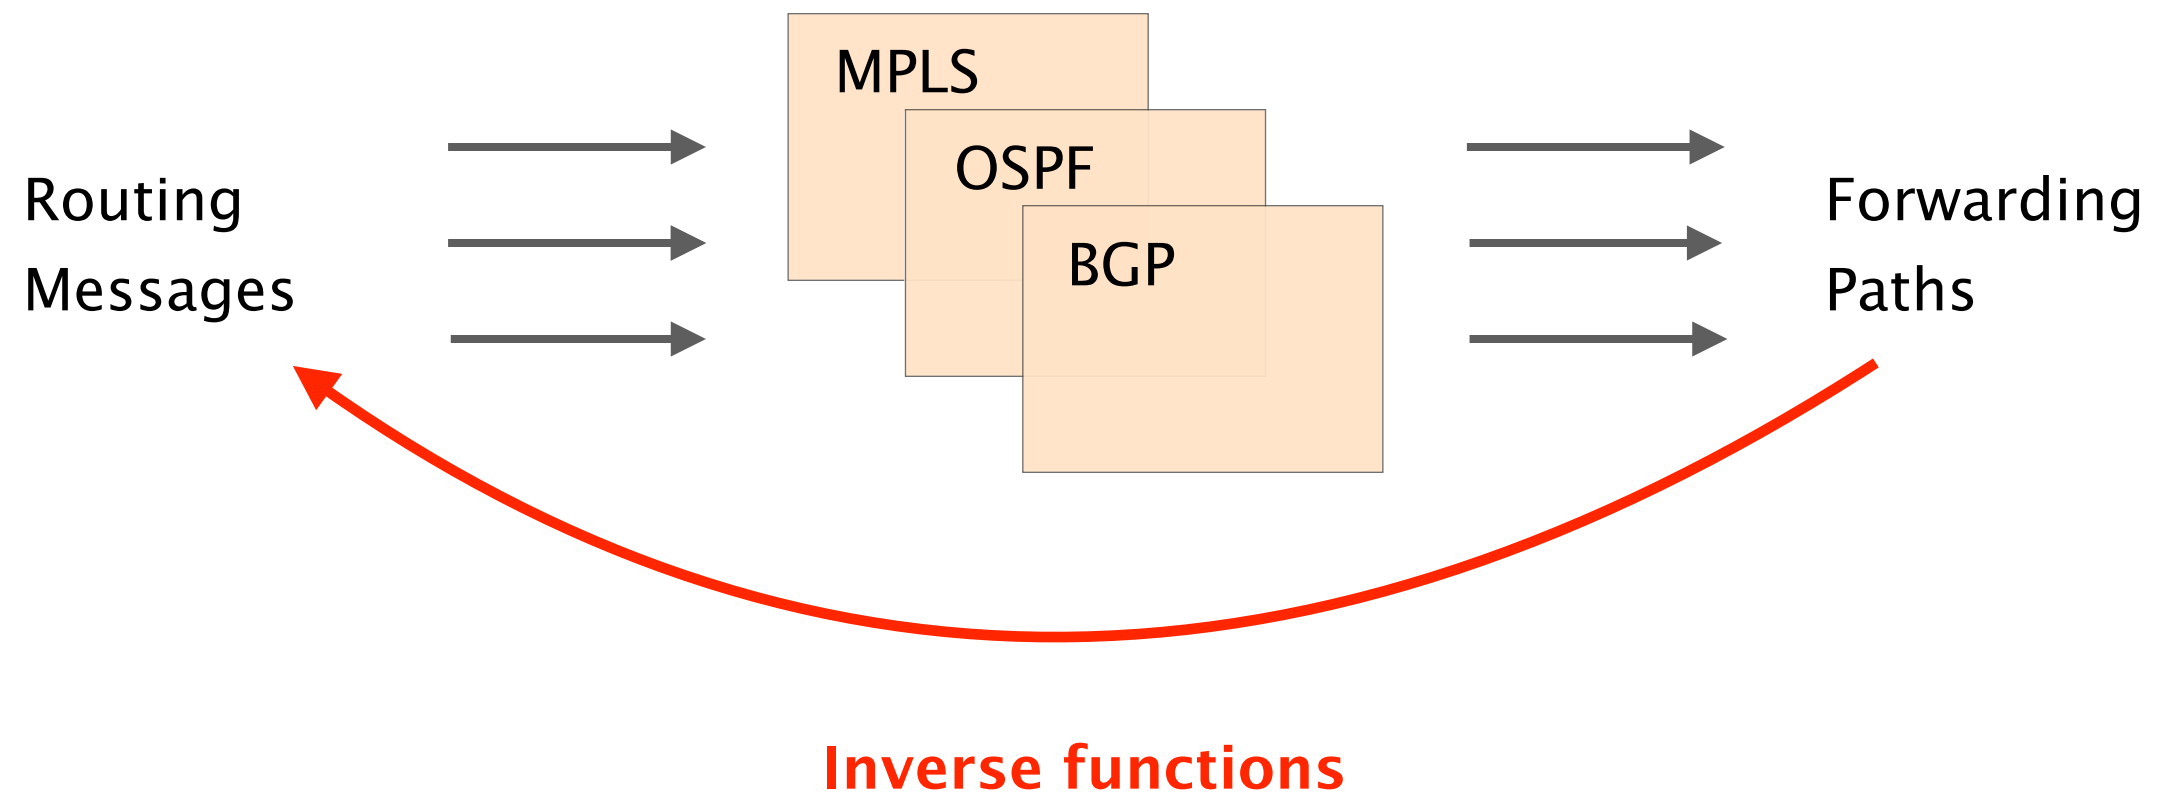
Functions are well known



Forwarding paths are also known,
from network-wide requirements



Given a forwarding path and a function,
can we automatically find the corresponding input?



The type of input to be computed depends on the routing protocol

	Type	Algorithm	Input
IGP	Link-State	Dijkstra	Network topology
BGP	Path-Vector	Decision Process	Received routes

On integrating Software-Defined Networking within existing routing systems



SDN-controlled routers
don't trash, recycle

2

SDN-controlled IGP
fine-grained traffic-engineering

SDN-controlled BGP
inter domain bonanza

Joint work with

Stefano Vissicchio, Olivier Bonaventure, Jennifer Rexford

Traffic Engineering techniques differ in terms of ease of use, functionality and support

IGP
(link reweight)

MPLS
(RSVP-TE)

SDN

signaling

expressiveness

device support

Traffic Engineering techniques differ in terms of ease of use, functionality and support

	IGP (link reweight)	MPLS (RSVP-TE)	SDN
signaling	no	yes	no
expressiveness	low shortest path	high	high
device support	excellent	good require MPLS	poor new hardware

In a SDN-controlled IGP, a controller presents a virtual topology to the routers to force them to use given paths

Given a set of forwarding paths,
augment an IGP topology with virtual

- nodes
- links and weights
- destinations

such that routers compute compatible paths

SDN-controlled IGP combines the benefits of each technique

SDN-controlled IGP

signaling

expressiveness

device support

SDN-controlled IGP combines the benefits of each technique

SDN-controlled IGP

signaling

no

expressiveness

high

device support

excellent

SDN-controlled IGP combines the benefits of each technique

SDN-controlled IGP

signaling

no

expressiveness

high

device support

excellent


SDN-controlled IGP enables fine-grained IP traffic control

SDN-controlled IGP enables to

- steer traffic on non-shortest paths
- forward along multiple paths (on a per-destination basis)
- provision backup paths


in a **centralized** manner, **on existing network**

Consider this network where a source sends traffic to 2 destinations




As congestion appears on the (C,D) link, operators might want to move away the orange flow to A

initial




desired




Moving only the orange flow to A is impossible with an IGP as both destinations are connected to D

initial




desired




*impossible to achieve by
reweighing the IGP links*

We can attract the orange flow from C by adding a virtual node announcing the orange destination





Traffic to V₁ is physically sent to A



Consider another network
with 2 sources and destinations





Consider another network
with 2 sources and destinations





The red and orange flows are limited to 100Mbps




If the two flows do not overlap all the time,
using 2 paths would enable each flow to use 200Mbps




If the two flows do not overlap all the time,
using 2 paths would enable each flow to use 200Mbps



Unfortunately, this is impossible to do with an IGP
as both destinations are connected to the same node




In contrast, SDN-controlled IGP enables
to create load-balanced paths on a per-destination basis




In contrast, SDN-controlled IGP enables
to create load-balanced paths on a per-destination basis

achieved ✓



virtual topology



A SDN-enabled IGP is powerful

Theorem	A SDN-enabled IGP can make the routers use any set of non-contradictory paths
---------	---

A SDN-enabled IGP is powerful

Theorem

A SDN-enabled IGP can make the routers use any set of **non-contradictory** paths

A SDN-enabled IGP is powerful

Theorem

A SDN-enabled IGP can make the routers use any set of **non-contradictory** paths

any path is loop-free

(*e.g.*, [s1, a, b, a, d] is not possible)


paths are consistent

(*e.g.* [s1, a, b, d] and

[s2, b, a, d] are inconsistent)

Given a physical topology and a set of path requirements, a linear program computes a virtual topology

physical
topology



+

paths requirements

[RA, RB, RC, RD]

ECMP(
[RX, RY, RZ],
[RX, RW, RZ])

[RI, RJ, RK], bckp
[RI, RL, RK]


⇒

Integer
Linear Program

Optimizer

⇒

virtual
topology



minimize topology size

SDN-enabled IGP is implementable in practice

SDN-enabled IGP requires to

- listen to the IGP traffic
simple, just establish an IGP adjacency
- inject fake IGP packets over an adjacency
effectively, executing a “controlled” IGP attack
- map virtual nodes to physical link
simple protocol change or use a few SDN-enabled devices

On integrating Software-Defined Networking within existing routing systems



SDN-controlled routers
don't trash, recycle

SDN-controlled IGP
fine-grained traffic-engineering

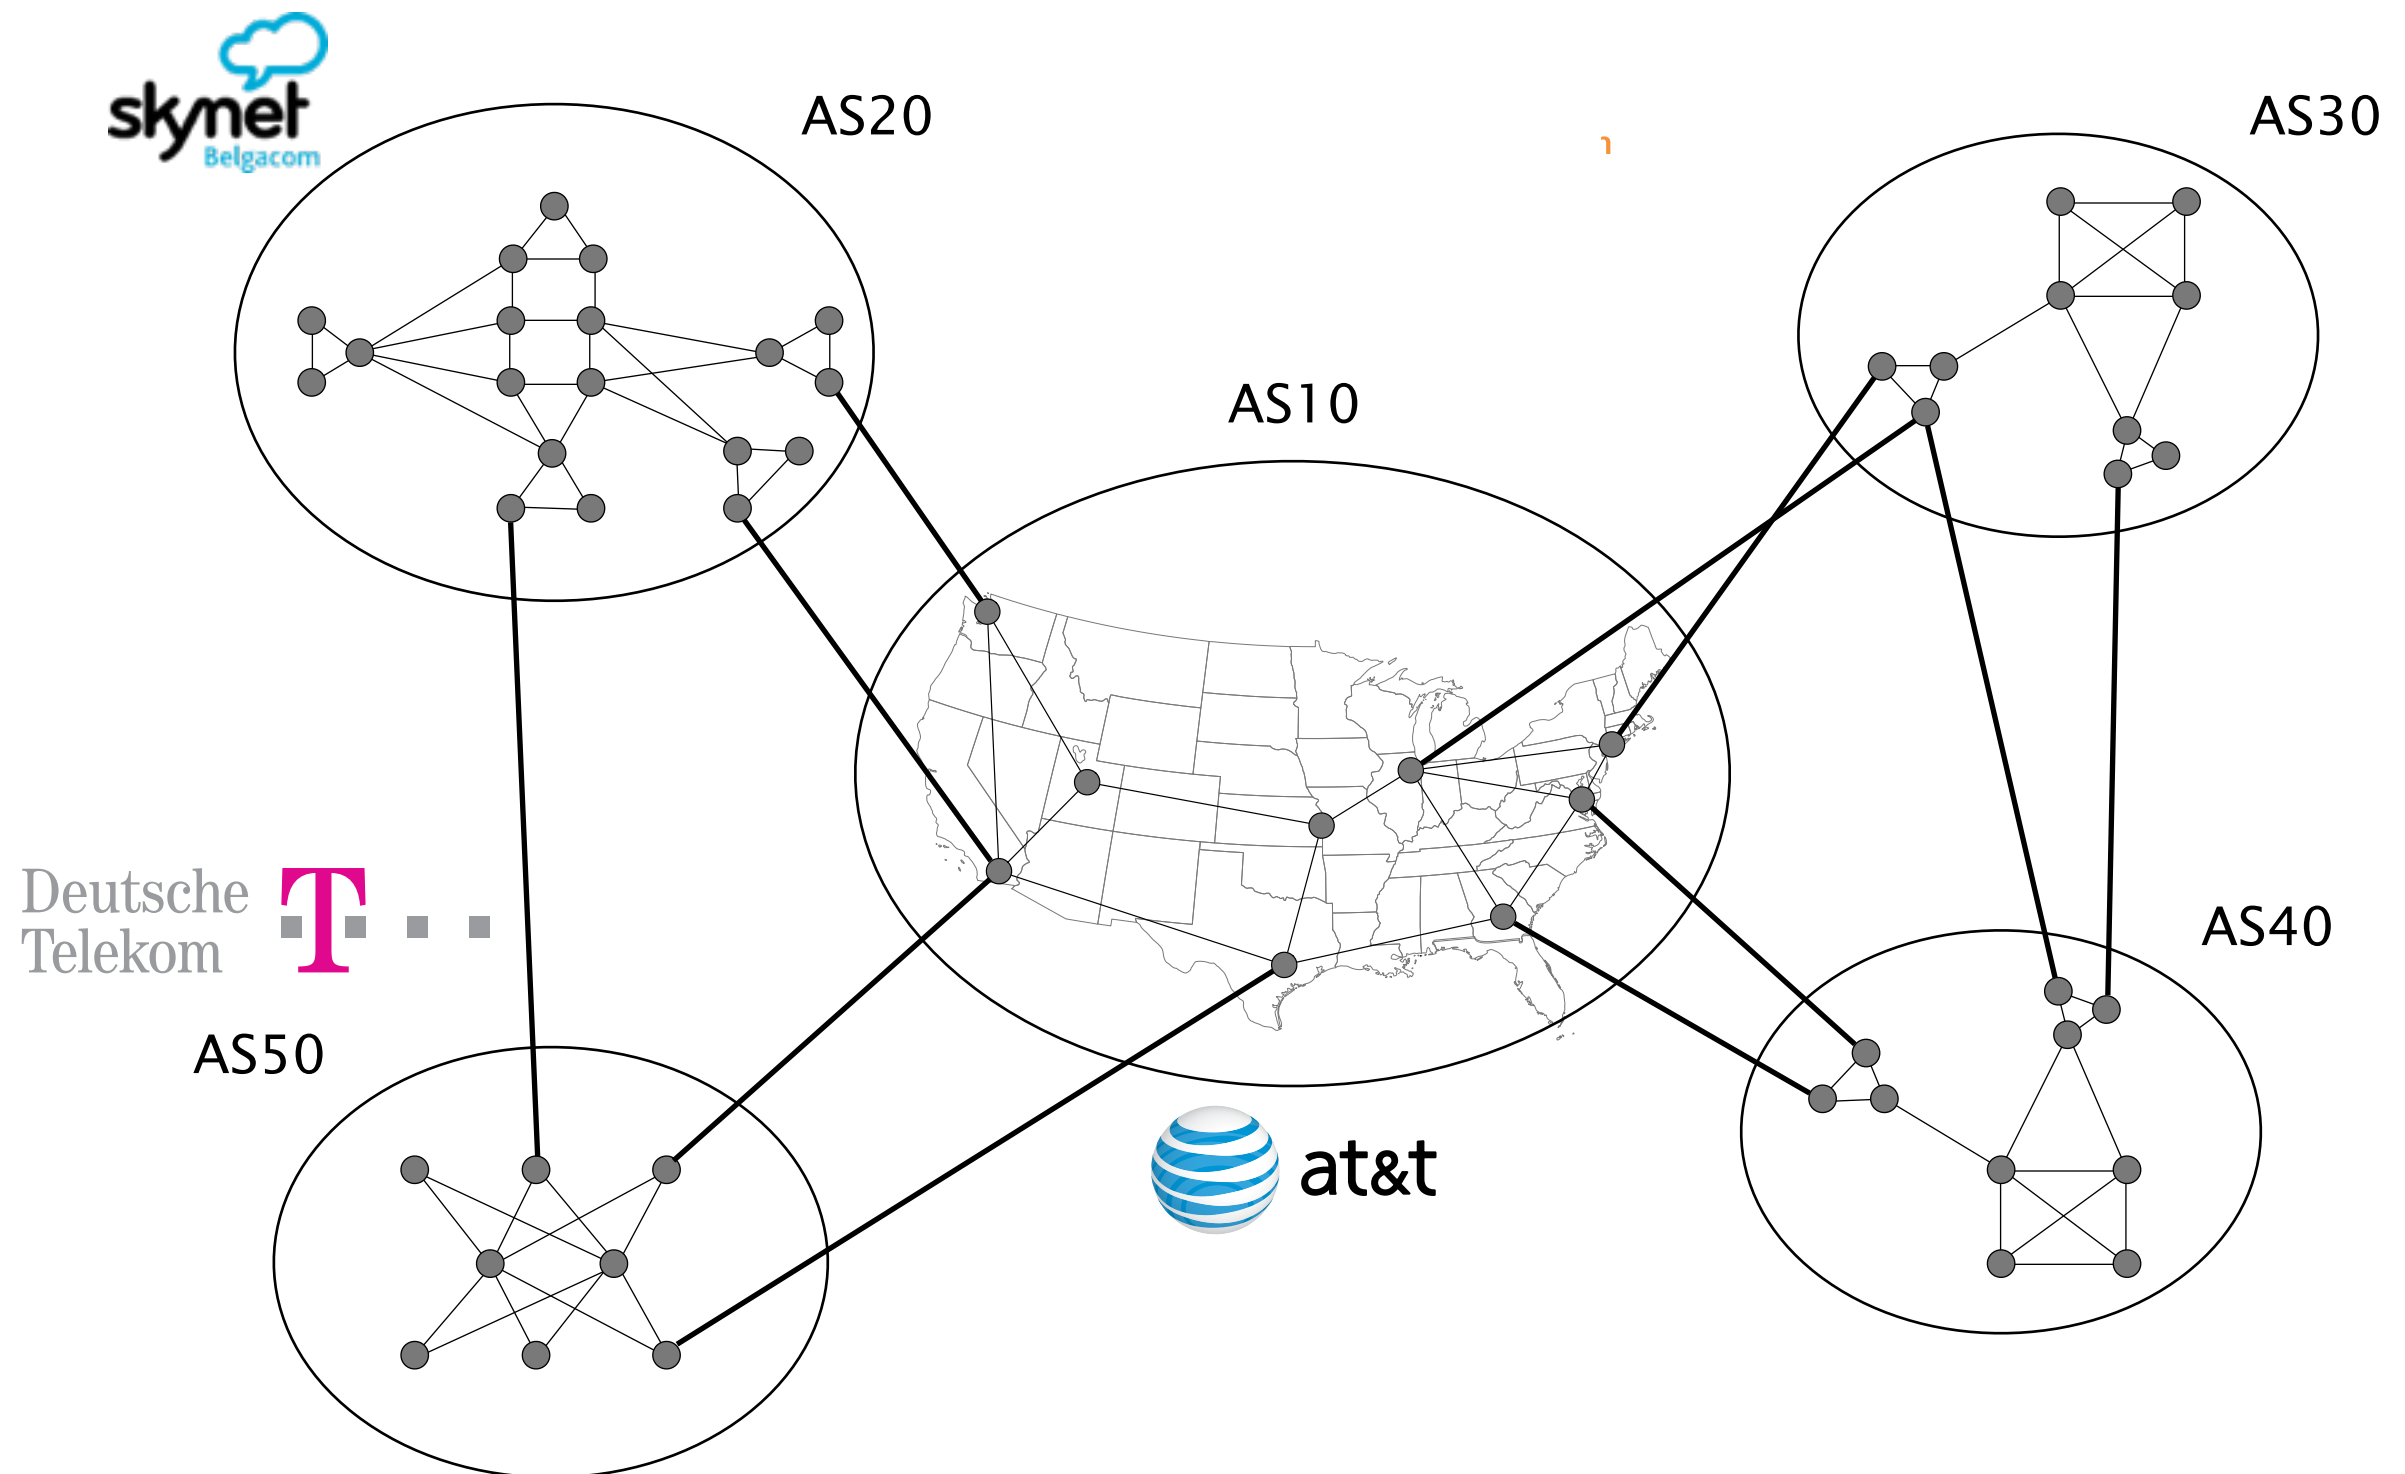
3

SDN-controlled BGP
inter domain bonanza

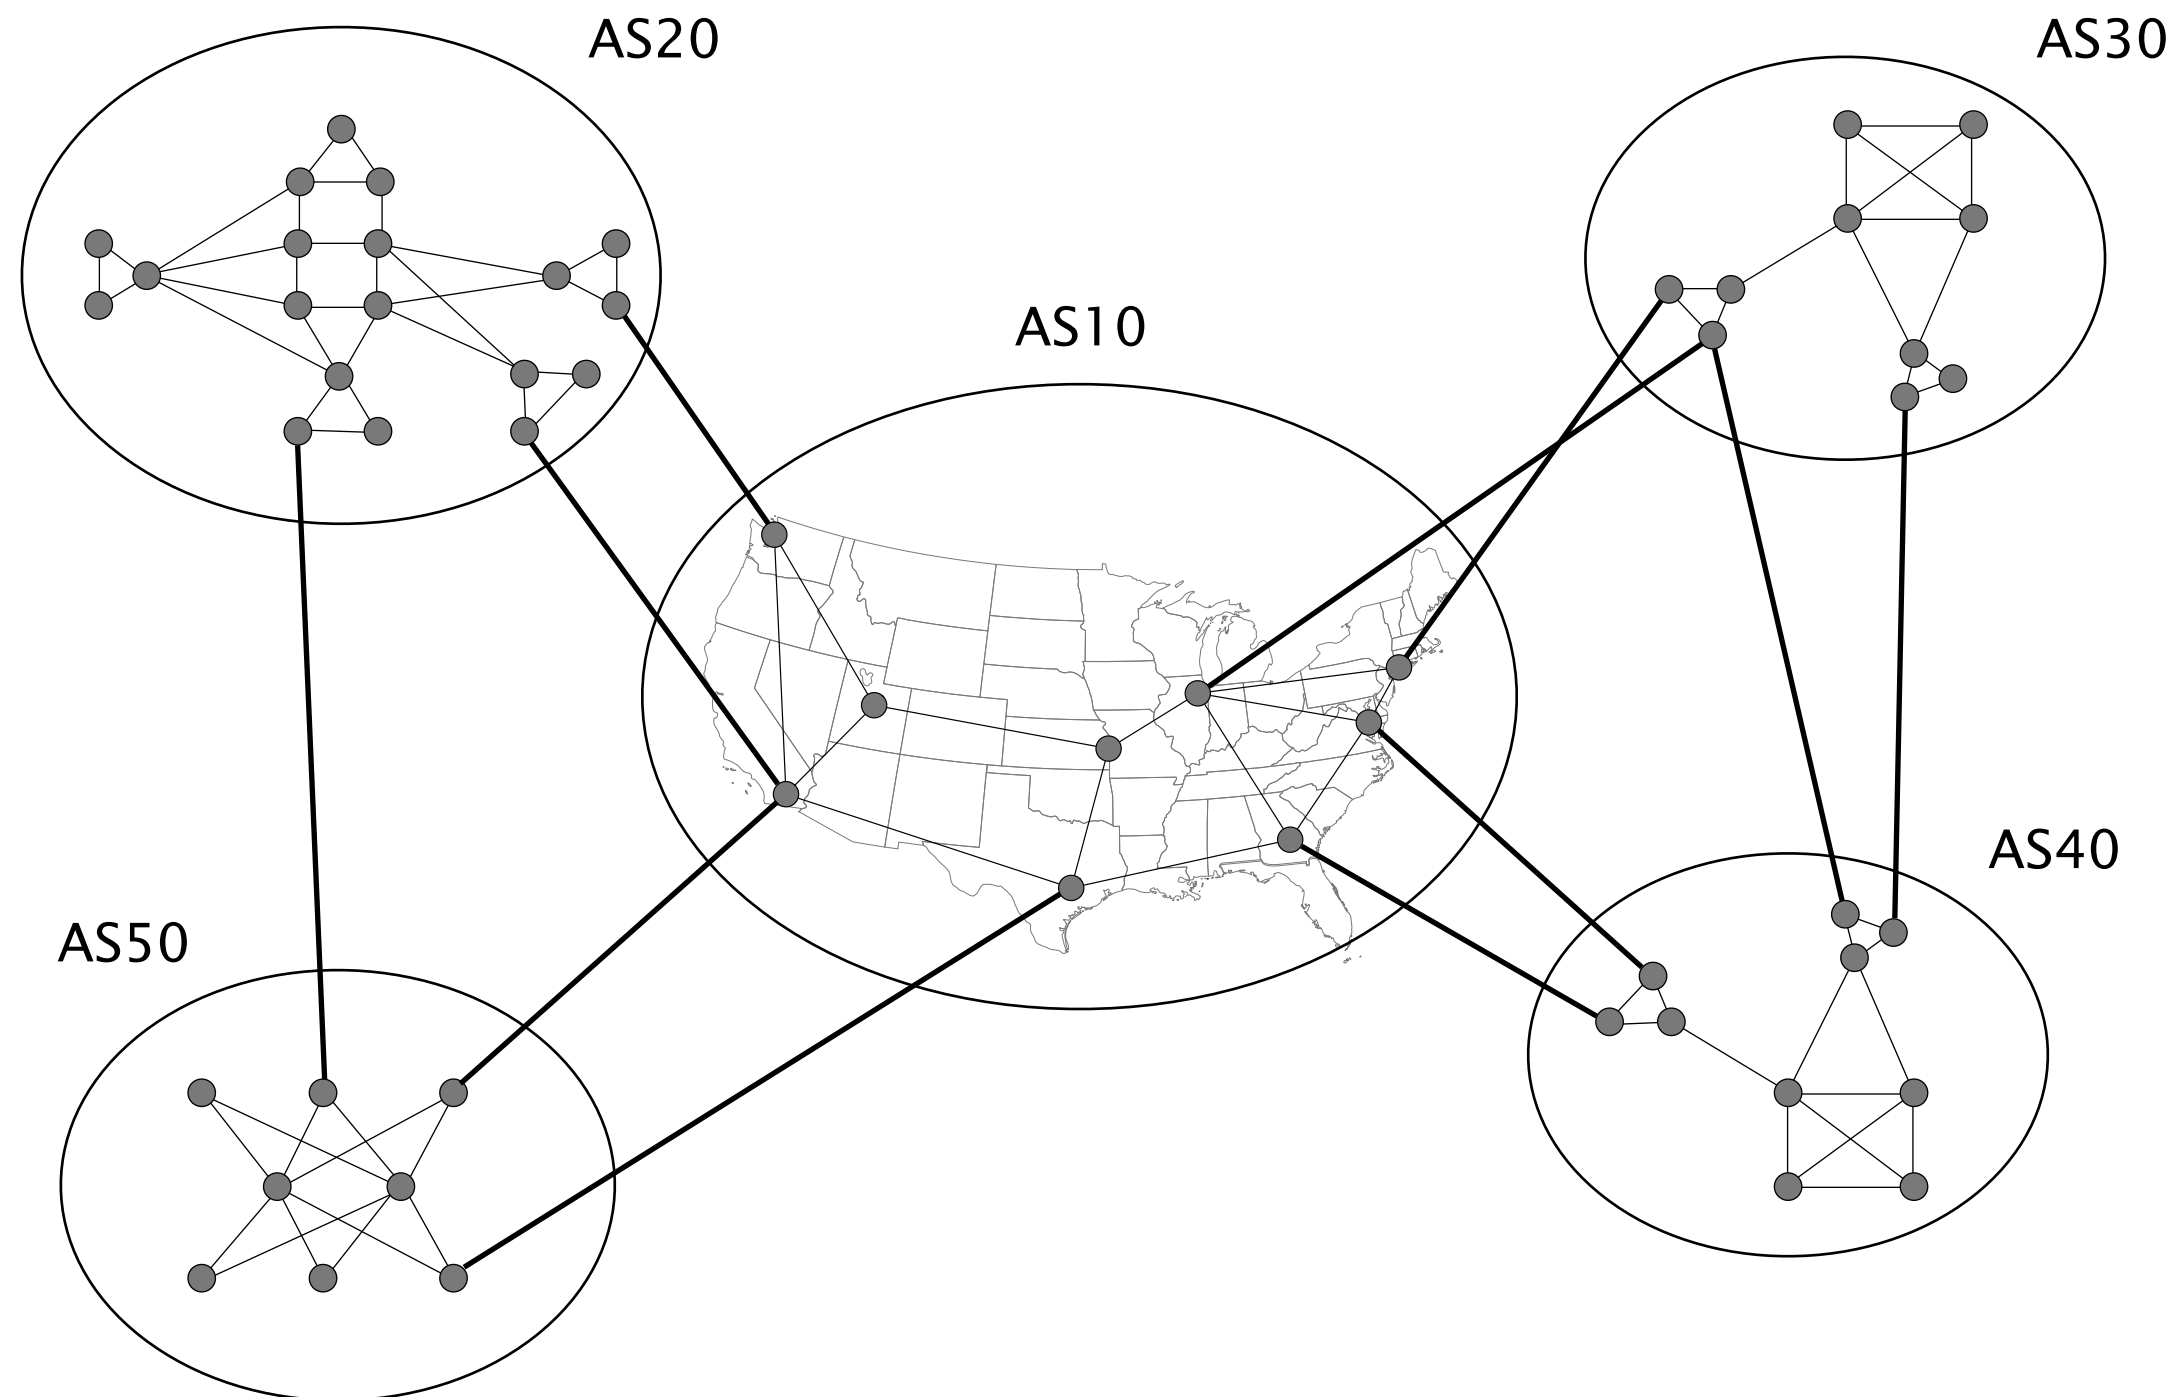
Joint work with

Arpit Gupta, Muhammad Shahbaz, Hyojoon Kim,
Russ Clark, Nick Feamster, Jennifer Rexford and Scott Shenker


BGP rules traffic forwarding across routing domains




BGP comes in two flavors



external BGP (eBGP) exchanges reachability information between Autonomous Systems (AS)



internal BGP (eBGP) distributes externally learned information within the AS



BGP can be (and is already) used as a centralized provisioning interface

Three examples of SDN-enabled BGP initiatives

- Route Control Platform [NSDI05]
- BGP Route Injection [LINX69]
- A BGP-Only SDN Controller for Large-Scale Data Centers [NANOG58]

Existing initiatives have focused on iBGP

iBGP

- Route Control Platform

iBGP

- BGP Route Injection

iBGP

- A BGP-Only SDN Controller for Large-Scale Data Centers

But managing external BGP is also painful
and would also benefit from SDN-like mechanisms

eBGP is

Inflexible (control-plane and data-plane)

BGP decision process and destination-based fwd

Non-deterministic

one can only “influence” remote decisions

Geographically-limited


one can only “do” something *where* it has an eBGP session

We combine BGP with SDN-enabled devices at Internet eXchange Points (IXP)


Augment the IXP data-plane with SDN capabilities
keeping default forwarding and routing behavior

Enable fine-grained inter domain policies
bringing new features while simplifying operations


We combine BGP with SDN-enabled devices at Internet eXchange Points (IXP)

- 
- Augment the IXP data-plane with SDN capabilities
keeping default forwarding and routing behavior
 - Enable fine-grained inter domain policies
bringing new features while simplifying operations
 - ... with **scalability** in mind
supporting the load of a large IXP


An IXP is a large L2 domain where participant routers exchange routes using BGP




To alleviate the need of establishing eBGP sessions, IXP often provides a Route Server (route multiplexer)




IP traffic is exchanged directly between participants, *i.e.* the IXP is forwarding transparent




With respect to a traditional IXP, SDN-enabled IXP (SDX)



With respect to a traditional IXP, SDN-enabled IXP (SDX) data-plane relies on SDN-capable devices



With respect to a traditional IXP, SDN-enabled IXP (SDX) control-plane relies on a SDN controller



SDX participants express their policies
in a high-level language built on top of Pyretic (*)

(*) <http://frenetic-lang.org/pyretic/>

SDX policies are composed of
a *pattern* and some *actions*

```
match ( Pattern ), then ( Actions )
```

Pattern selects packets based on any header fields,

Pattern

```
match ( eth_type  
        vlan_id  
        srcmac  
        dstmac , && , || ), then ( Actions )  
        protocol  
        dstip  
        tos  
        srcip  
        srcport  
        dstport
```

Pattern selects packets based on any header fields,
while actions forward or modify the selected packets

match (**Pattern**), then (**Actions**)

drop
forward
rewrite

Each participant writes her policies *independently* and transmits them to the controller

Participant A's policy:

```
match(dstip=ipA.1), fwd(A1)  
match(dstip=ipA.2), fwd(A2)
```

Participant B's policy:

```
match(dstip=ipC), fwd(C)  
match(dstip=ipA), fwd(A)  
match(dstip=ipB), fwd(B)
```

```
match(dstip=ipC), fwd(C)
```

Participant C's policy

SDN controller



Given the participant policies, the controller compiles them to SDN forwarding entries

Ensuring isolation

Resolving policies conflict

Ensuring scalability

Given the participant policies, the controller compiles them to SDN forwarding entries

Ensuring isolation



Each participant controls one virtual switch

connected to participants it can communicate with

Resolving policies conflict


Ensuring scalability

Given the participant policies, the controller compiles them to SDN forwarding entries

Ensuring isolation

Resolving policies conflict

Ensuring scalability



Participant policies are sequentially composed


in an order that respects business relationships

Given the participant policies, the controller compiles them to SDN forwarding entries

Ensuring isolation

Resolving policies conflict

Ensuring scalability




only install the minimum
required in the data plane

leverage the existing BGP
control plane for the rest


The edge routers, sitting next to the fabric,
are tailored to match on numerous IP prefixes

not FIB-constrained

FIB constrained




Edge router




SDN switch


We consider routers FIB as the first stage of a multi-stage FIB




Routers FIB match on the destination prefix
and set a tag accordingly



The SDN FIB matches on the tag,
not on the IP prefixes



How do we provision tag entries in a router, and what are these tags?



We use BGP as a provisioning interface
and the L2 address of the BGP NH as label

When a BGP router receives a route, it

- runs the decision process
- resolves the IP next-hop to a L2 next-hop (if best)
- installs a FIB entry directing traffic to the L2 next-hop


We use BGP as a provisioning interface
and the L2 address of the BGP NH as label

When a BGP router receives a route, it

- runs the decision process
- resolves the IP next-hop to a L2 next-hop (if best)
- installs a FIB entry directing traffic to the L2 next-hop

We can tweak the IP/L2 next-hop and use it as a tag

Let's walk through the compilation of a simple inbound TE policy



AS B's SDX policy

<code>match(dstip=0*)</code>	<code>fwd(B1)</code>
<code>match(dstip=1*)</code>	<code>fwd(B2)</code>

A, B and C are all connected to the SDX

The policy is first divided
in match and forward actions

match(dstip=0*)	fwd(B1)
match(dstip=1*)	fwd(B2)

The policy is first divided
in match and forward actions


match(dstip=0*)

match(dstip=1*)

fwd(B1)


fwd(B2)

A virtual IP/MAC next-hop is associated to each distinct forwarding actions



The SDX controller provisions two data plane rules matching the destination MAC

fwd(B1)	——	(NH1, MAC1)
fwd(B2)	——	(NH2, MAC2)




Forwarding rules

match(dst:MAC1), fwd(B1)
match(dst:MAC2), fwd(B2)

The SDX controller rewrite the BGP NH of B's routes according to the match part of the policy

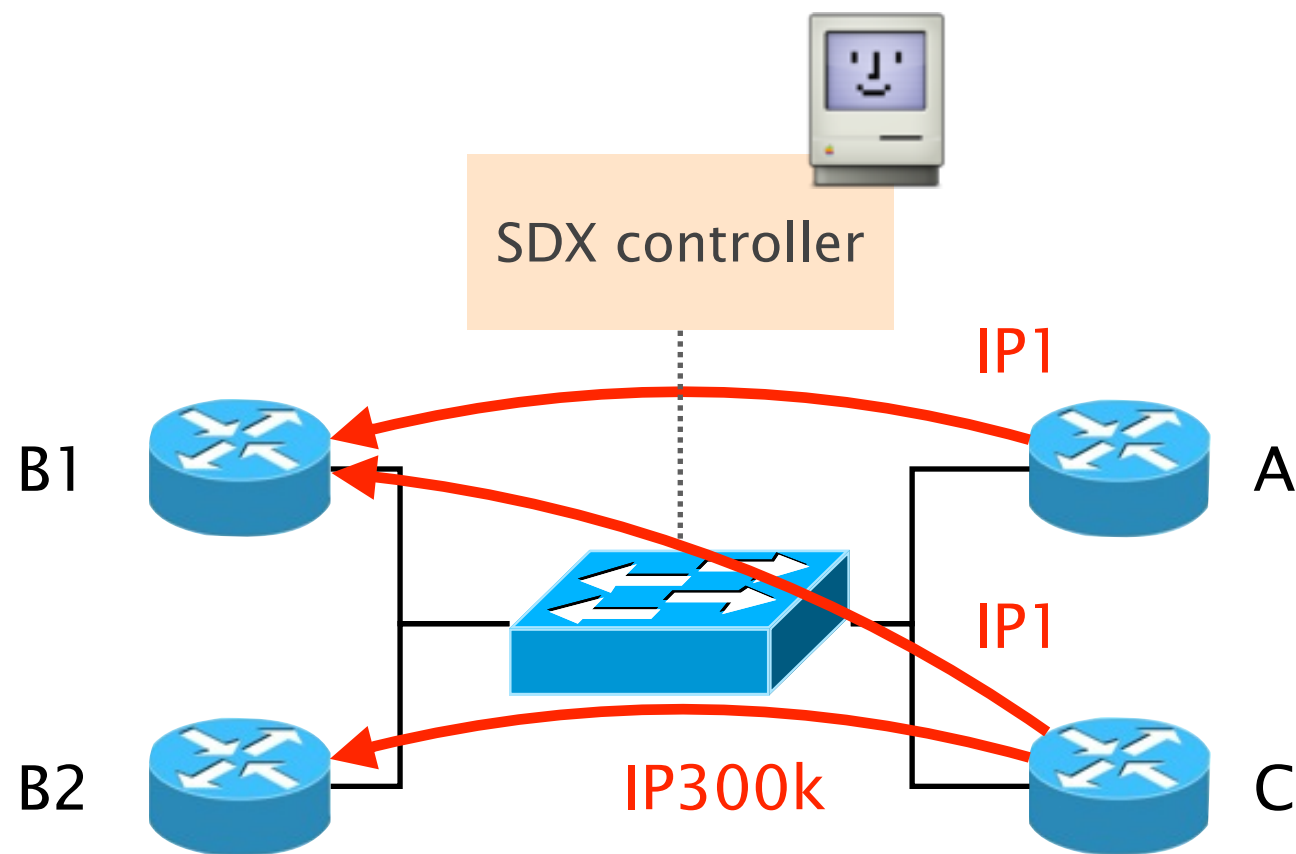
`match(dstip=0*)` — (NH1, MAC1)
`match(dstip=1*)` — (NH2, MAC2)



BGP routes sent to A & C

<i>prefix</i>	<i>NH</i>
p1	NH1
...	...
p250k	NH1
<hr/>	
p250k+1	NH2
...	...
p250k	NH2

Traffic from A and C is splitted on B1 and B2 according to B's policy, **with only 2 data-plane rules**



What else does SDX enable that was
hard or **impossible** to do before?

SDX enables a wide range of novel applications

security

Prevent/block policy violation

Prevent participants communication

forwarding optimization

Middlebox traffic steering

Traffic offloading

Inbound Traffic Engineering

Fast convergence

peering

Application-specific peering

remote-control

Upstream blocking of DoS attacks

Influence BGP path selection

Wide-area load balancing

SDX enables a wide range of novel applications

security

Prevent/block policy violation

Prevent participants communication

forwarding optimization

Middlebox traffic steering

Traffic offloading

Inbound Traffic Engineering

Fast convergence

peering

Application-specific peering

remote-control

Upstream blocking of DoS attacks

Influence BGP path selection

Wide-area load balancing

DNS-based wide-area load balancing has several limitations


High TTL values lead to slow recovery when a device fails
due to caching by local DNS servers and browsers

Low TTL values lead to higher delay for DNS resolution
due to cache misses


Load-balancing is not based on the client IP address
but on the DNS resolver IP address (*e.g.*, 8.8.8.8)

SDX enable direct and quick control
of traffic redirection

Let's consider a CDN C that provides one service at two Data Centers (DC)




C assigns one IP prefixes per DC




C assigns one IP address identifying the service

 **192.0.10.1** ——— 192.0.10.0/24 is a **service** prefix



C announces the service prefix at the IXP



C directs the service traffic to the appropriate DC based on the client's IP address

```
match(dstip=192.0.10.1) then
  (match(srcip=0.0.0.0/1) then
    mod(dstip=192.0.1.161)) and      // forward to DC1
  (match(srcip=128.0.0.0/1) then
    mod(dstip=192.0.2.139))          // forward to DC2
```

SDX enables direct and quick control of traffic redirection

SDX-based load-balancing is

- fast no DNS caching problem
- flexible use of any load-balancing algorithm
- efficient based on the actual client IP address

On integrating Software-Defined Networking within existing routing systems



SDN-controlled routers
don't trash, recycle

SDN-controlled IGP
fine-grained traffic engineering

SDN-controlled BGP
inter domain bonanza

SDN-controlled routing enables to realize parts
of the SDN promises today, on an existing network

Facilitate a complete transition to SDN

provide *one interface to rule them all*

Simplify the controller implementation

most of the hard work is done by the routers

Maintain operators' mental model

same good old protocols running, easier troubleshooting

On integrating Software-Defined Networking within existing routing systems



Laurent Vanbever

www.vanbever.eu

Stanford University

November 13 2013